

Online Vehicle Booking Market Segmentation Analysis

LADI JEEVAN SAI

ML INTERNSHIP

TASK 2



Dataset Link:-

https://drive.google.com/file/d/1_hwV8vWkVMFwB1XiiEoigKpCGqeVJ15g/view?usp=share_link

Project Link:-

[GitHub - ladijeevansai/Online-vehicle-booking-market-segmentation](https://github.com/ladijeevansai/Online-vehicle-booking-market-segmentation)

Dataset description

I had collected the data from kaggle ,which consists of 14 attributes described below and consists of 131662 rows.

The attributes are described as follows:

- Trip_ID: A unique identifier for each trip.
- Trip_Distance: The distance in units unspecified that the customer has requested for the trip.
- Type_of_Cab: A categorical variable indicating the category of the cab requested by the customer.
- Customer_Since_Months: A numerical variable indicating the number of months the customer has been using cab services; 0 indicates the current month.
- Life_Style_Index: A proprietary index created by Sigma Cabs that indicates the lifestyle of the customer based on their behavior.
- Confidence_Life_Style_Index: A categorical variable indicating the level of confidence in the Life_Style_Index mentioned above.
- Destination_Type: Sigma Cabs divides any destination into one of 14 categories.
- Customer_Rating: The average lifetime rating of the customer until the present time.
- Cancellation_Last_1Month: The number of trips canceled by the customer in the last month.
- Var1, Var2, and Var3: Continuous variables masked by the company that can be used for modeling purposes.
- Gender: A categorical variable indicating the gender of the customer.
- Surge_Pricing_Type: The target variable indicating the surge pricing type; it can take on 3 different values.

```
[ ] print(df.shape)
(131662, 14)

[ ] df.head(10)
```

	Trip_ID	Trip_Distance	Type_of_Cab	Customer_Since_Months	Life_Style_Index	Confidence_Life_Style_Index	Destination_Type	Customer_Rating	Cancellation_Last_1Month	Var1	Var2	Var3
0	T0005689460	6.77	B	1.0	2.42769		A	3.90500	0	40.0	46	60
1	T0005689461	29.47	B	10.0	2.78245		B	3.45000	0	38.0	56	78
2	T0005689464	41.58	NaN	10.0	NaN		NaN	3.50125	2	NaN	56	77
3	T0005689465	61.56	C	10.0	NaN		NaN	3.45375	0	NaN	52	74
4	T0005689467	54.95	C	10.0	3.03453		B	3.40250	4	51.0	49	102
5	T0005689469	19.06	E	10.0	NaN		NaN	2.59750	1	72.0	63	91
6	T0005689470	29.72	E	10.0	2.83958		C	2.97500	1	83.0	50	75
7	T0005689472	18.44	B	2.0	2.81871		B	3.58250	0	103.0	46	63
8	T0005689473	106.80	C	3.0	NaN		NaN	3.14625	0	NaN	58	92
9	T0005689474	107.19	D	5.0	3.04467		B	2.44375	1	NaN	58	83

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 131662 entries, 0 to 131661
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Trip_ID                               131662 non-null  object
1   Trip_Distance                         131662 non-null  float64
2   Type_of_Cab                           111452 non-null  object
3   Customer_Since_Months                 125742 non-null  float64
4   Life_Style_Index                      111469 non-null  float64
5   Confidence_Life_Style_Index           111469 non-null  object
6   Destination_Type                      131662 non-null  object
7   Customer_Rating                       131662 non-null  float64
8   Cancellation_Last_1Month              131662 non-null  int64
9   Var1                                  60632 non-null   float64
10  Var2                                  131662 non-null  int64
11  Var3                                  131662 non-null  int64
12  Gender                                131662 non-null  object
13  Surge_Pricing_Type                    131662 non-null  int64
dtypes: float64(5), int64(4), object(5)
memory usage: 14.1+ MB
```

Data Preprocessing

In this step , initially i had replaced all the null values with mean if it is float aor integer data type and replaced by mode if it is categorical attribute

```
[ ] df.isna().sum()

Trip_ID                0
Trip_Distance          0
Type_of_Cab            20210
Customer_Since_Months  5920
Life_Style_Index       20193
Confidence_Life_Style_Index  20193
Destination_Type        0
Customer_Rating         0
Cancellation_Last_1Month  0
Var1                   71030
Var2                   0
Var3                   0
Gender                 0
Surge_Pricing_Type     0
dtype: int64
```

```

for col in df.columns:
    if df[col].dtype == 'object':
        # For object type columns, replace null values with the mode
        df[col].fillna(df[col].mode()[0], inplace=True)
    else:
        # For numeric type columns, replace null values with the mean
        df[col].fillna(df[col].mean(), inplace=True)

df.head(10)

```

	Trip_ID	Trip_Distance	Type_of_Cab	Customer_Since_Months	Life_Style_Index	Confidence_Life_Style_Index	Destination_Type	Customer_Rating	Cancellation_Last_1Month	Var1	Var2
0	T0005689460	6.77	B	1.0	2.427690		A	3.90500	0	40.000000	46
1	T0005689461	29.47	B	10.0	2.782450		B	3.45000	0	38.000000	56
2	T0005689464	41.58	B	10.0	2.802064		B	3.50125	2	64.202698	56
3	T0005689465	61.56	C	10.0	2.802064		B	3.45375	0	64.202698	52
4	T0005689467	54.95	C	10.0	3.034530		B	3.40250	4	51.000000	49
5	T0005689469	19.06	E	10.0	2.802064		B	2.59750	1	72.000000	63
6	T0005689470	29.72	E	10.0	2.839580		C	2.97500	1	83.000000	50
7	T0005689472	18.44	B	2.0	2.818710		B	3.58250	0	103.000000	46
8	T0005689473	106.80	C	3.0	2.802064		B	3.14625	0	64.202698	58
9	T0005689474	107.19	D	5.0	3.044670		B	2.44375	1	64.202698	58

```

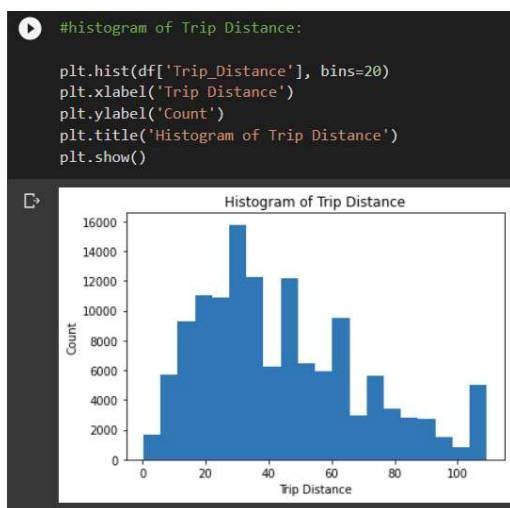
df.isna().sum()

```

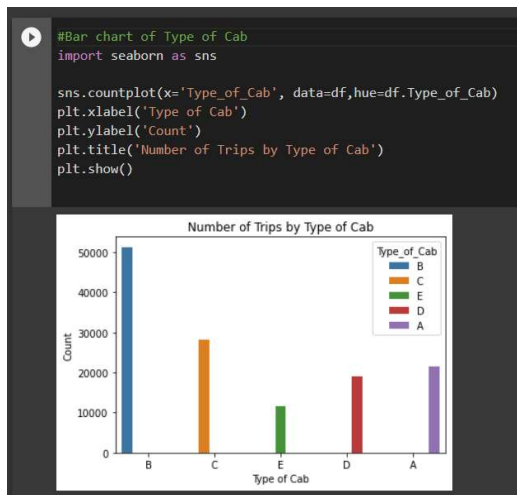
Trip_ID	0
Trip_Distance	0
Type_of_Cab	0
Customer_Since_Months	0
Life_Style_Index	0
Confidence_Life_Style_Index	0
Destination_Type	0
Customer_Rating	0
Cancellation_Last_1Month	0
Var1	0
Var2	0
Var3	0
Gender	0
Surge_Pricing_Type	0
dtype: int64	

Data visualization:-

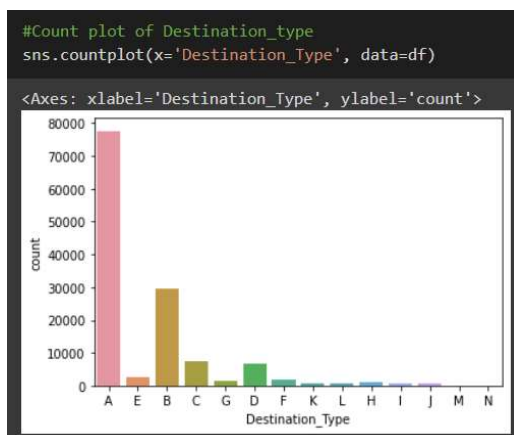
Using matplotlib and seaborn libraries i had tried to visualize several attributes in the dataset



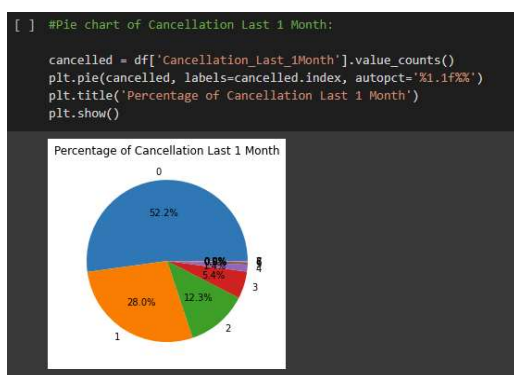
This plot shows about the frequency distribution of the Trip distance attribute.



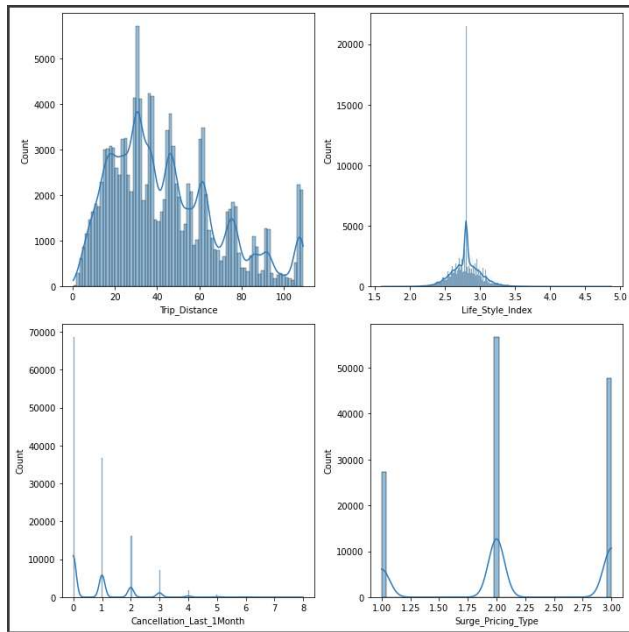
This plot tells that the sigma cabs have 'B' type cabs the most and 'E' type the least.



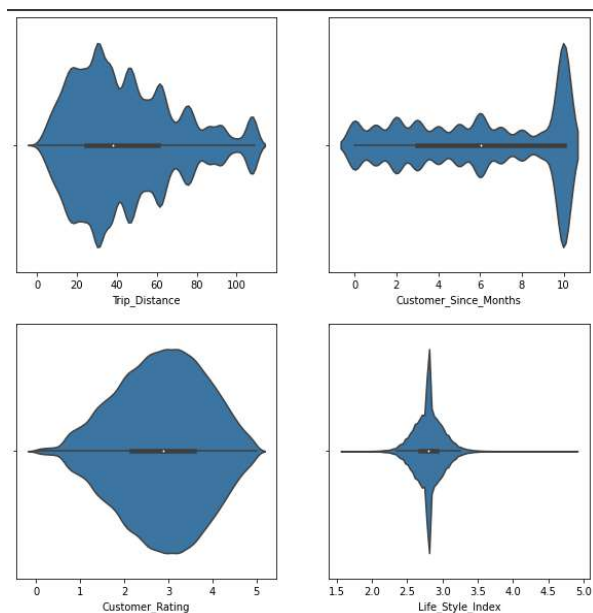
The plot shows that most people prefer Destination type A, while types F, G, K, L, H, I, J, M, and N are very unpopular. This suggests Sigma cabs can supply more cabs at Destinations A and B, while having very few or no cabs at Destinations F, G, K, L, H, I, J, M, N.

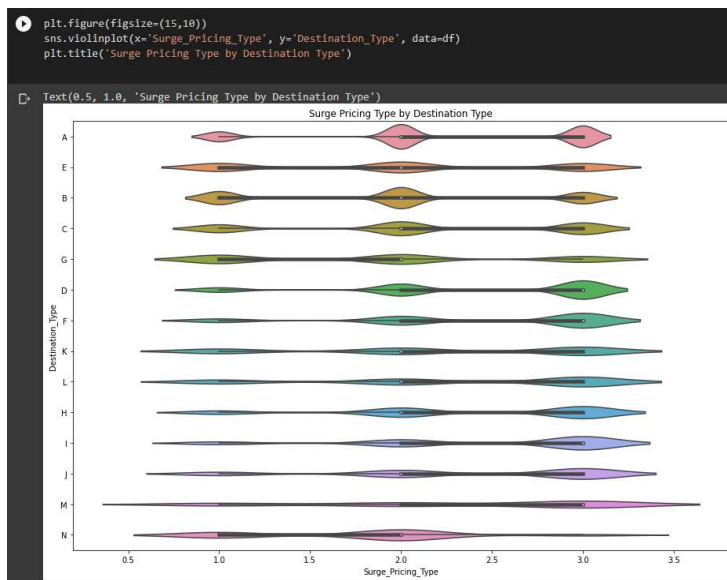


This shows us the percentage distribution of the ride cancellations in last month



Based on the plots, it can be inferred that the majority of people who opted for cabs have chosen it for a good distance, with the distance ranging between 20-60. Additionally, the lifestyle index of most people falls between 2.5-3.25, with an average of 2.75. Furthermore, the third plot shows that the cancellation rate is not very high and ranges mostly between 1-3, with around 10,000 people never canceling their ride.

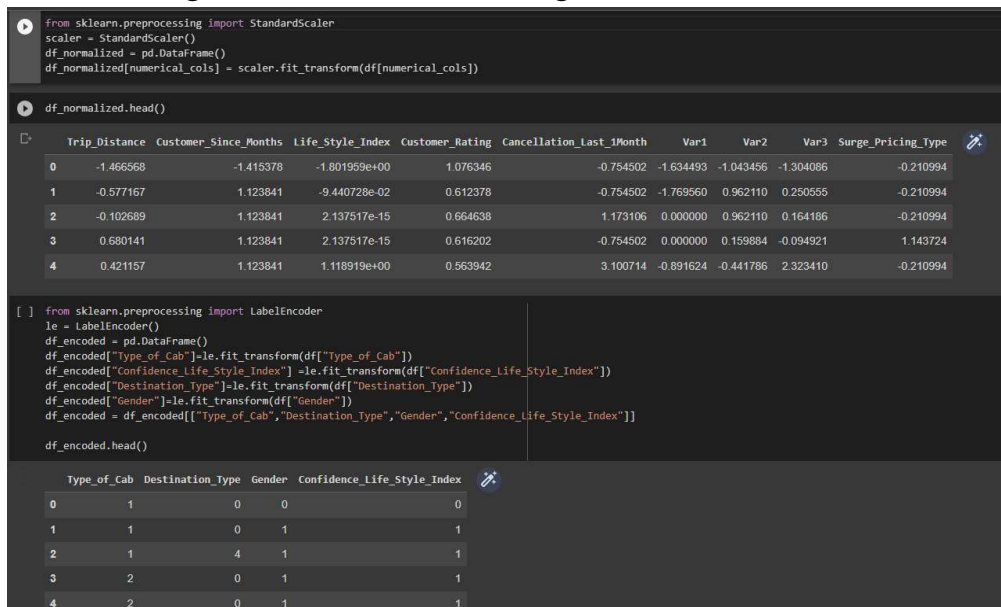




The above visualizations are in the form of violin plots, which display the complete distribution of several variables. Specifically, they show the distribution of trip distance, customer service months, customer rating, and lifestyle index. Additionally, there is a violin plot depicting the relationship between surge pricing and destination type, which allows us to identify locations with higher surge rates and provides valuable insights into various aspects of the dataset, including the range and distribution of key variables, as well as patterns and trends in the data.

By analyzing these plots, we can gain a better understanding about the relation between the dataset and make more informed decisions based on this information.

Normalizing and one hot encoding the data



```

] # Concatenated vertically
df_final = pd.concat([df_normalized, df_encoded], axis=1)

] df_final.head()

```

	Trip_Distance	Customer_Since_Months	Life_Style_Index	Customer_Rating	Cancellation_Last_1Month	Var1	Var2	Var3	Surge_Pricing_Type	Type_of_Cab	Destination_Type	Gender	Confidence_Life_Style_Index
0	-1.466568	-1.415378	-1.801959e+00	1.076346	-0.754502	-1.634493	-1.043456	-1.304086	-0.210994	1	0	0	0
1	-0.577167	1.123841	-9.440728e-02	0.612378	-0.754502	-1.769560	0.962110	0.250555	-0.210994	1	0	1	1
2	-0.102689	1.123841	2.137517e-15	0.664638	1.173106	0.000000	0.962110	0.164186	-0.210994	1	4	1	1
3	0.680141	1.123841	2.137517e-15	0.616202	-0.754502	0.000000	0.159884	-0.094921	1.143724	2	0	1	1
4	0.421157	1.123841	1.118919e+00	0.563942	3.100714	-0.891624	-0.441786	2.323410	-0.210994	2	0	1	1

In the above code i had followed few steps to preprocess the data:

1. Separated the numerical and categorical attributes into separate arrays.
2. Normalized the numerical attributes (integer and float) using StandardScaler, which is a common method of normalization that scales the data to have a mean of 0 and a standard deviation of 1.
3. Encoded your categorical attributes into numerical values, which is a common method of transforming non-numeric data into a format that can be used by machine learning algorithms.
4. Merged the two dataframes back together to create a single, uniform, and normalized dataset.

This preprocessing pipeline can help improve the performance of your machine learning models by reducing the impact of differences in scale and format between different types of attributes. Normalizing numerical attributes can make them more comparable, while encoding categorical attributes can make them more amenable to machine learning algorithms.

Removing Outliers and dropping unnecessary attributes:-

```

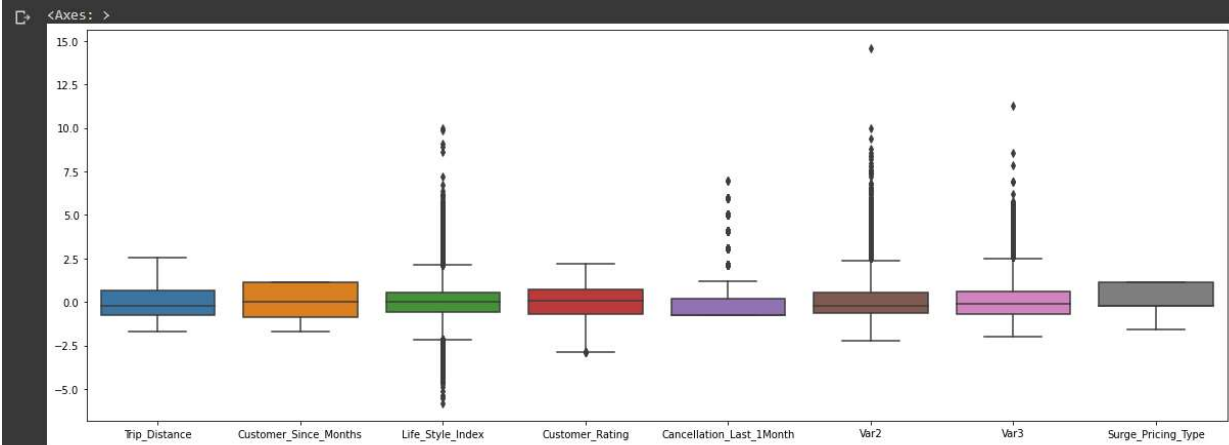
] #correlation matrix
corr_matrix = df_final.corr()
corr_matrix

```

	Trip_Distance	Customer_Since_Months	Life_Style_Index	Customer_Rating	Cancellation_Last_1Month	Var1	Var2	Var3	Surge_Pricing_Type	Type_of_Cab	Destination_Type	Gender	Confidence_Life_Style_Index
Trip_Distance	1.000000	0.114340	0.468367	-0.054654	-0.007696	-0.030662	0.200456	0.231706	0.135928	0.067647	-0.171064	0.002173	
Customer_Since_Months	0.114340	1.000000	0.119154	-0.049001	-0.006187	-0.006700	0.041766	0.110830	0.027202	0.009622	-0.054814	0.001327	
Life_Style_Index	0.468367	0.119154	1.000000	0.189173	0.068176	-0.055873	0.215921	0.303296	-0.073682	-0.045736	-0.022695	0.000653	
Customer_Rating	-0.054654	-0.049001	0.189173	1.000000	0.003595	-0.007400	-0.302968	-0.227531	-0.155279	-0.091945	0.133342	0.000468	
Cancellation_Last_1Month	-0.007696	-0.006187	0.068176	0.003595	1.000000	0.007550	0.095830	0.128686	0.185646	0.092795	0.047085	0.002307	
Var1	-0.030662	-0.006700	-0.055873	-0.007400	0.007550	1.000000	-0.031712	-0.041235	-0.026302	-0.014182	0.001769	0.000354	
Var2	0.200456	0.041766	0.215921	-0.302968	0.095830	-0.031712	1.000000	0.683437	0.003437	-0.001347	-0.093573	0.003183	
Var3	0.231706	0.110830	0.303296	-0.227531	0.128686	-0.041235	0.683437	1.000000	-0.039309	-0.022756	-0.079056	0.003520	
Surge_Pricing_Type	0.135928	0.027202	-0.073682	-0.155279	0.185646	-0.026302	0.003437	-0.039309	1.000000	0.503093	0.012152	0.000872	
Type_of_Cab	0.067647	0.009622	-0.045736	-0.091945	0.092795	-0.014182	-0.001347	-0.022756	0.503093	1.000000	0.003836	0.001103	
Destination_Type	-0.171064	-0.054814	-0.022695	0.133342	0.047085	0.001769	-0.093573	-0.079056	0.012152	0.003836	1.000000	-0.003070	
Gender	0.002173	0.001327	0.000653	0.000468	0.002307	0.000354	0.003183	0.003520	0.000872	0.001103	-0.003070	1.000000	
Confidence_Life_Style_Index	0.224798	0.035793	0.119489	-0.061443	0.028031	-0.005231	0.044439	0.055519	0.115344	0.080250	-0.039013	0.005123	


```
[ ] df_numeric = df_final[numerical_cols]
df_numeric = df_numeric.drop('Var1', axis=1)
```

```
# Draw boxplots for all numerical columns
fig = plt.figure(figsize=(20,7))
sns.boxplot(data=df_numeric)
```

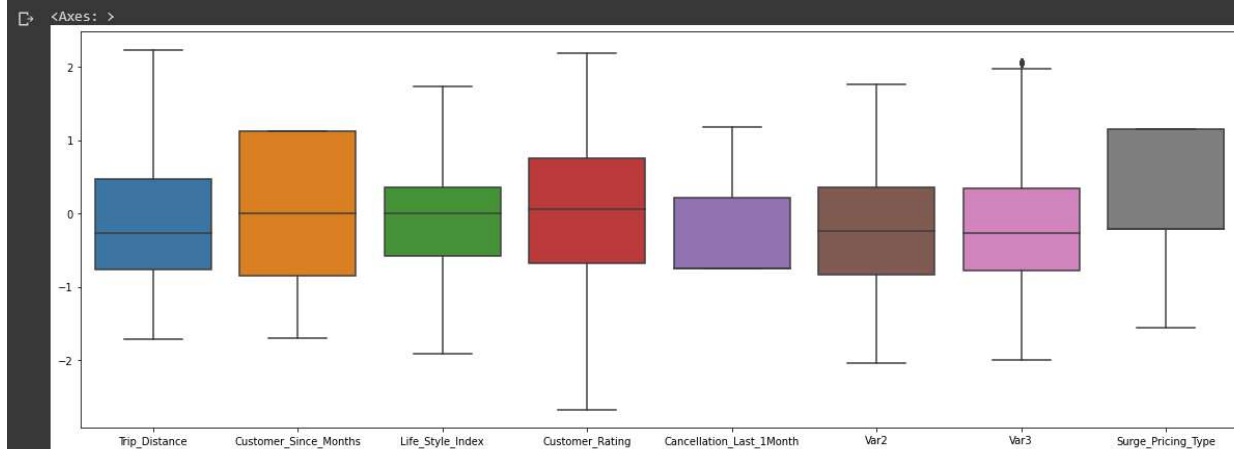


Removing outliers

df_N_NO= df_numeric data frame with no outliers

```
[ ] Q1 = df_numeric.quantile(0.25)
Q3 = df_numeric.quantile(0.75)
IQR = Q3 - Q1
df_N_NO = df_numeric[~((df_numeric < (Q1 - 1.5 * IQR)) | (df_numeric > (Q3 + 1.5 * IQR))).any(axis=1)]
```

```
# Draw boxplots for all numerical columns
fig = plt.figure(figsize=(20,7))
sns.boxplot(data=df_N_NO)
```



```
print(df_numeric.shape)
```

```
(131662, 8)
```

```
print(df_N_NO.shape)
```

```
(101994, 8)
```

- I have created a correlation matrix to analyze the relationships between attributes, and have decided to exclude the attribute var1 as it has the lowest correlation with most of the other variables.
- Then I had plotted box plots to examine the numerical attributes and found that there were a significant number of outliers. After removing the outliers, I had observed that the number of rows in the data set reduced by approximately 30,000, and then confirmed the absence of outliers by plotting the box plots again.

Applying PCA for dimensionality reduction

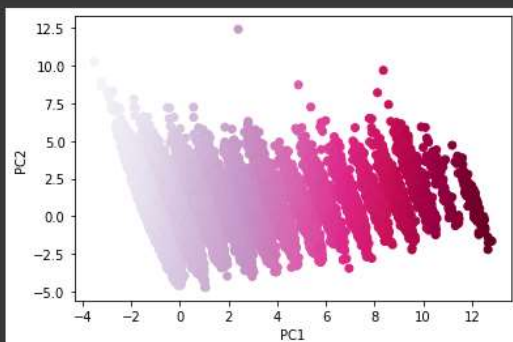
1. Two Principal components

```
[ ] from sklearn.decomposition import PCA

# Apply PCA
pca = PCA(n_components=2)
principal_components = pca.fit_transform(df_final)

# Create a new DataFrame with the principal components
df_pca = pd.DataFrame(data=principal_components, columns=['PC1', 'PC2'])

# Visualize the first two principal components
plt.scatter(df_pca['PC1'], df_pca['PC2'], c=df_pca['PC1'], cmap='PuRd')
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.show()
```



```
df_pca.head()
```

	PC1	PC2
0	-0.175221	-3.283565
1	-1.031841	-0.050944
2	2.821884	1.252343
3	-1.087845	0.258809
4	-1.237786	2.026508

2. Three Principal components

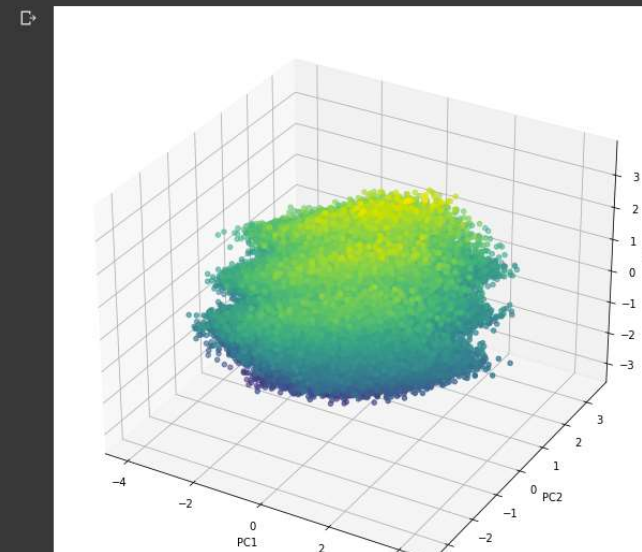
```
from sklearn.decomposition import PCA

# Apply PCA
pca2 = PCA(n_components=3)
principal_components2 = pca2.fit_transform(df_N0)

# Create a new DataFrame with the principal components
df_pca2 = pd.DataFrame(data=principal_components2, columns=['PC1', 'PC2', 'PC3'])
```

Visualize the three principal components

```
fig = plt.figure(figsize=(10, 10))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(df_pca2['PC1'], df_pca2['PC2'], df_pca2['PC3'], c=df_pca2['PC3'], cmap='viridis')
ax.set_xlabel('PC1')
ax.set_ylabel('PC2')
ax.set_zlabel('PC3')
plt.show()
```



I had used PCA to reduce the dimensionality of the dataset. As there were 14 attributes and 131662 rows. By applying PCA, we can reduce the number of attributes (columns) while retaining most of the variability in the data. This can be helpful for clustering, and classification. Additionally, reducing the number of attributes can also help in reducing the computational cost of algorithms.

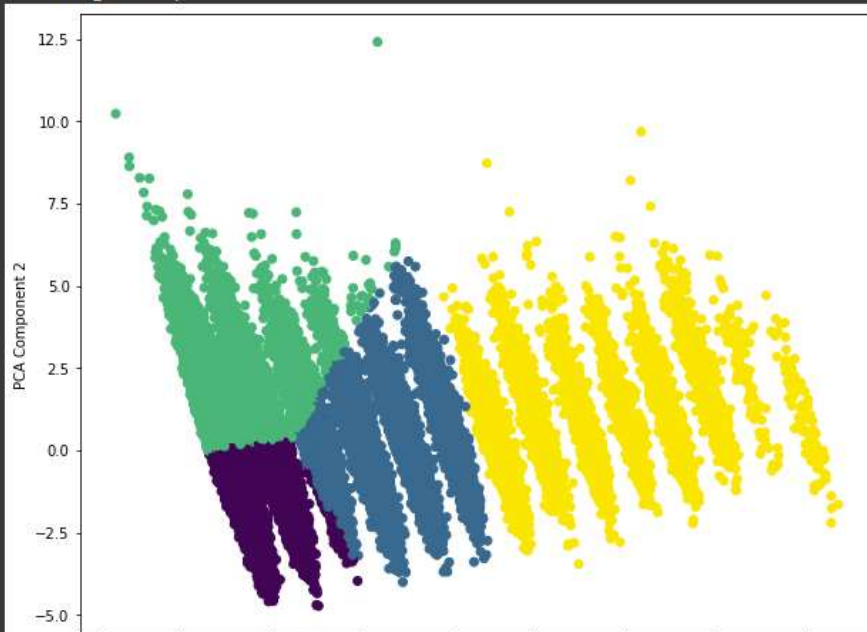
Segmentation using the data upon which PCA is applied

```
# Perform KMeans clustering
kmeans = KMeans(n_clusters=4, random_state=0)
clusters = kmeans.fit_predict(df_pca)

# Visualize the clusters
import matplotlib.pyplot as plt

fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111)
ax.scatter(df_pca.values[:, 0], df_pca.values[:, 1], c=clusters, cmap='viridis')
ax.set_xlabel('PCA Component 1')
ax.set_ylabel('PCA Component 2')
plt.show()
```

/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: warnings.warn()



```
# Calculate metrics
from sklearn.metrics import silhouette_score
inertia = kmeans.inertia_
silhouette = silhouette_score(df_pca, clusters)
```

```
print('Inertia:', inertia)
print('Silhouette score:', silhouette)
```

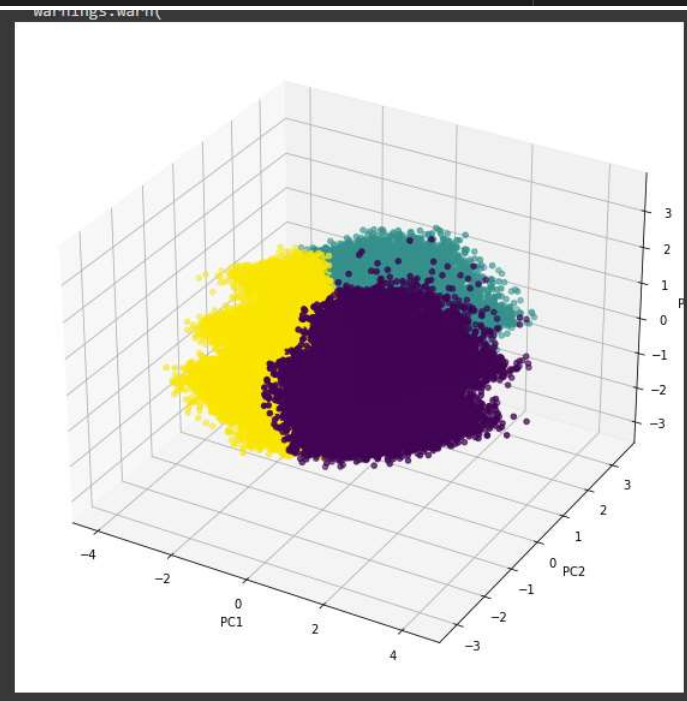
```
Inertia: 192276.1693878618
Silhouette score: 0.4461127405408211
```

```
# Perform KMeans clustering
kmeans2 = KMeans(n_clusters=3, random_state=0)
clusters2 = kmeans2.fit_predict(df_pca2)
```

```
# Visualize the clusters
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
```

```
#Visualizing the plot for Customer_Rating Cancellation_Last_1Month Vari
```

```
# Visualize the first two principal components
fig = plt.figure(figsize=(10, 10))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(df_pca2.values[:, 0], df_pca2.values[:, 1], df_pca2.values[:, 2], c=clusters2, cmap='viridis')
ax.set_xlabel('PC1')
ax.set_ylabel('PC2')
ax.set_zlabel('PC3')
plt.show()
```



```
# Calculating metrics
from sklearn.metrics import silhouette_score
inertia = kmeans.inertia_
silhouette = silhouette_score(df_pca2, clusters2)

print('Inertia:', inertia)
print('Silhouette score:', silhouette)

Inertia: 14835.907519295968
Silhouette score: 0.23688338553475774
```

Use of K Means clustering for market segmentation

KMeans clustering is a popular unsupervised learning algorithm used for market segmentation. The primary goal of market segmentation is to divide customers into groups or clusters with similar characteristics so that companies can create targeted marketing strategies and improve customer experience.

KMeans clustering algorithm can be used for market segmentation on the above data by grouping customers with similar characteristics such as trip distance, type of cab, customer rating, etc. into clusters. The algorithm aims to minimize the distance between the data points within each cluster while maximizing the distance between the clusters. The result is a set of clusters that represent the different segments of customers.

Once the clusters are created, companies can tailor their marketing efforts to each segment, creating targeted campaigns that address the unique needs and preferences of each group. This approach can lead to increased customer loyalty, higher customer satisfaction, and ultimately, increased revenue for the company.

Overall, KMeans clustering can be a powerful tool for market segmentation, allowing companies to gain a deeper understanding of their customers and create more effective marketing strategies.

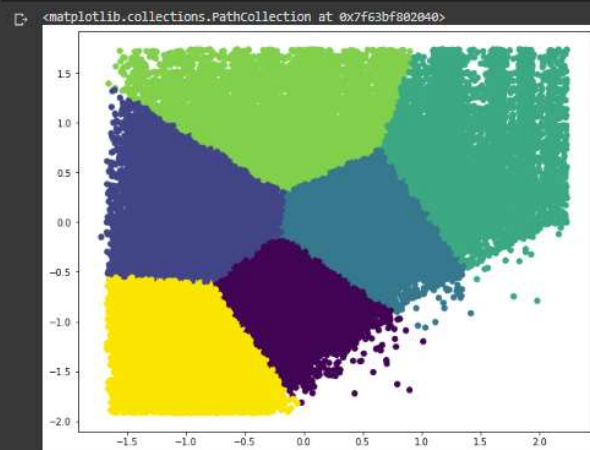
Segmentation using the data upon which PCA is applied

k means clustering for two columns

```
from sklearn.cluster import KMeans
kmc = KMeans(n_clusters=6,n_init=1)
model_kmc = kmc.fit(df_kmc)
model_kmc.labels_

array([5, 1, 2, ..., 4, 4, 4], dtype=int32)

fig = plt.figure(figsize=(10, 8))
plt.scatter(df_N_NO["Trip_Distance"],df_N_NO.Life_Style_Index,c=model_kmc.labels_)
```



```
# Calculating metrics
from sklearn.metrics import silhouette_score
inertia = kmeans.inertia_
silhouette = silhouette_score(df_kmc,model_kmc.labels_)

print('Inertia:', inertia)
print('Silhouette score:', silhouette)

Inertia: 14835.907519295968
Silhouette score: 0.3733336007588392
```

k means clustering for 3 columns

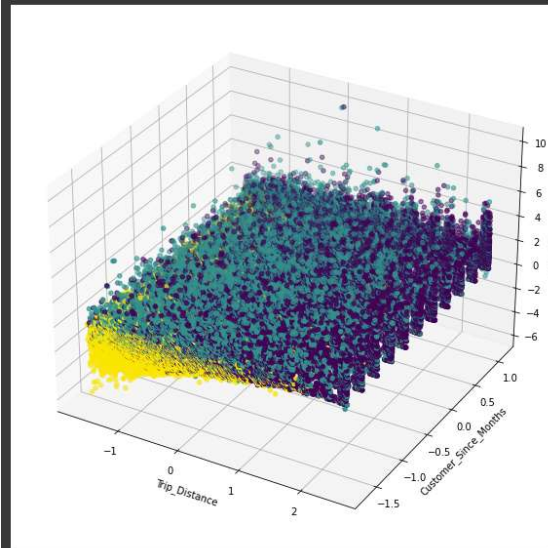
```
[ ] from sklearn.cluster import KMeans
    from sklearn.preprocessing import StandardScaler

    # Drop non-numeric columns
    #df_numeric = df_N_NO[numerical_cols]

    # Perform KMeans clustering
    kmeans = KMeans(n_clusters=3, random_state=0)
    clusters = kmeans.fit_predict(df_numeric)
```



```
fig = plt.figure(figsize=(10, 10))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(df_numeric.values[:, 0], df_numeric.values[:, 1], df_numeric.values[:, 2], c=clusters, cmap='viridis')
ax.set_xlabel(numerical_cols[0])
ax.set_ylabel(numerical_cols[1])
ax.set_zlabel(numerical_cols[2])
plt.show()
```



Conclusion:-

Based on the analysis of the dataset consisting of 131,662 rows with 14 attributes , several insights can be drawn. Firstly, there are outliers present in the dataset , which are be removed to improve the accuracy of the model. Secondly, PCA is used to reduce the dimensionality of the dataset and identify the most important variables. Then by using the elbow method, it became possible to identify three as the optimal number of clusters for KMeans clustering. So these can be used for market segmentation to target customers based on their trip distance, type of cab requested, customer rating, lifestyle index, confidence in lifestyle index, destination type, cancellation rate, and surge pricing type. Overall, the use of PCA and KMeans clustering can help the companies to identify patterns and clusters in the data, which can be used to make informed decisions in marketing and business strategies.